


When is specialization worth its cache footprint?


Specialist adapters buy quality; shared-prefix reuse and adapter capacity decide whether that quality is serveable.

A. Cache namespace mechanism




B. Real overlap trend

1.5B vLLM: 5 overlap levels x 3 seeds; 7B shown as anchors




C. Quality/goodput frontier

expanded source-backed Qwen2.5-7B eval



D. Adapter capacity boundary

startup feasibility at 4096 context



Decision rule: specialize when quality gain exceeds SLO loss, cache fragmentation, and capacity cost.